

# Multimedia Systems

## CS-538 Advanced Operating Systems

(Slides include materials from *Operating System Concepts*, 7<sup>th</sup> ed., by Silberschatz, Galvin, & Gagne and from *Modern Operating Systems*, 2<sup>nd</sup> ed., by Tanenbaum)

# Outline

- What is multi-media?
- Requirements and challenges for audio and video in computer systems
  - CD devices and formats
- Systems for multimedia
- Compression and bandwidth
- If time:–
  - Processor and disk scheduling
  - Network streaming and management

Silberschatz, Chapter 20  
Tanenbaum, Chapter 7

# What do we mean by “multimedia”

- Audio and video within a computer system
  - CD's & DVD's
  - Computer hard drive
  - Interactive
- Live broadcast & web casts
  - Radio stations, Network TV, Webcams, Skype, ...
- Video on demand
  - Pause, fast forward, reverse, etc.
- Interactive activities involving audio, video, images
  - Meetings and presentations with 2-way audio
  - Teleconferencing
- Handheld devices
  - iPod, MP-3 players; personal video; mobile phones, ...
- ...

# Multimedia Data and Delivery

- Stored in file system like ordinary data.
- Must be accessed with specific timing requirements.
- E.g., video *must* be displayed at 24-30 frames per second.
  - System must guarantee timely delivery
- *Continuous-media data*
  - Data with specific rate requirements

# Multimedia Characteristics

- Multimedia files can be *very* large.
- Continuous media data may require *very* high data rates.
- Multimedia applications are usually sensitive to timing delays during playback.
- *Note:* human ear is more sensitive to *jitter* in audio than eye is to *jitter* in video!

# Requirements

- “Smooth” audio and video
  - Deterioration in quality  $\Rightarrow$  jerky playback
- Multiple concurrent streams
  - Video & multimedia servers
  - TiVo, etc.
- Wide range of network bandwidths
- Audio/video while PC is doing something else

# System and OS Challenges

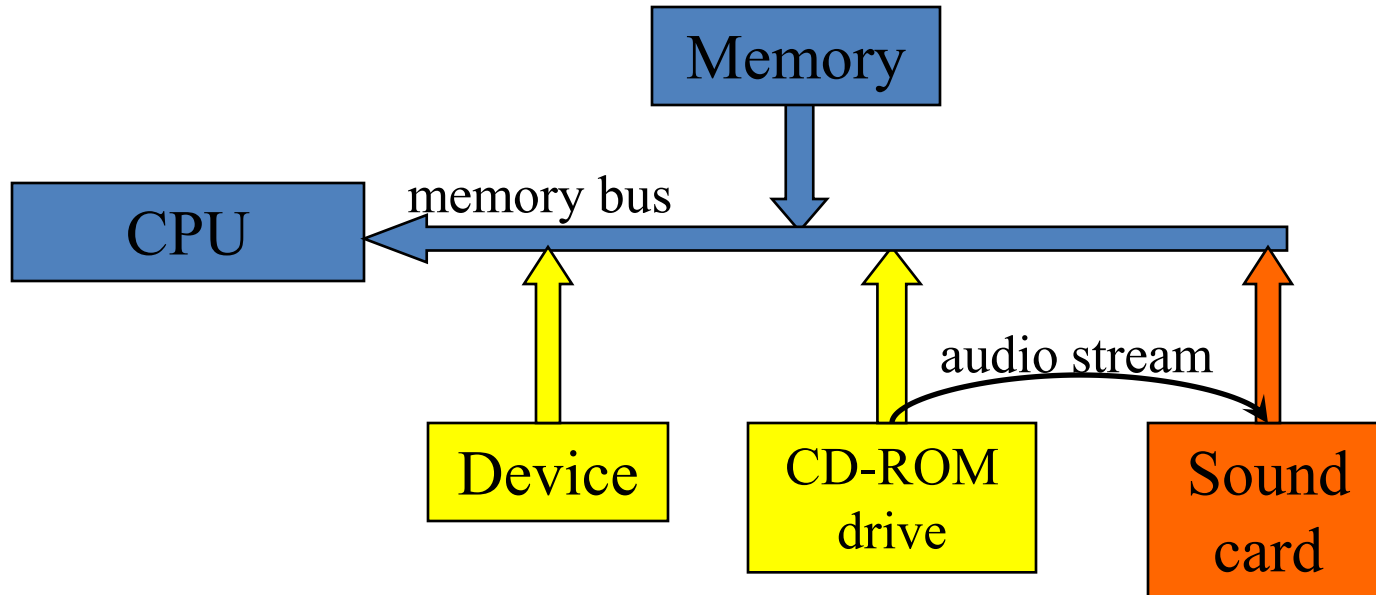
- Bandwidths and Compression
- Jitter
- Processor Scheduling
- Disk Scheduling
- Network Streaming

# Some System Architectures

- Simple:
  - Data paths for audio/video that are separate from computational data paths
- Modern
  - Fast system bus, CPU, devices
- Video server
  - Disk farm and multiple streams

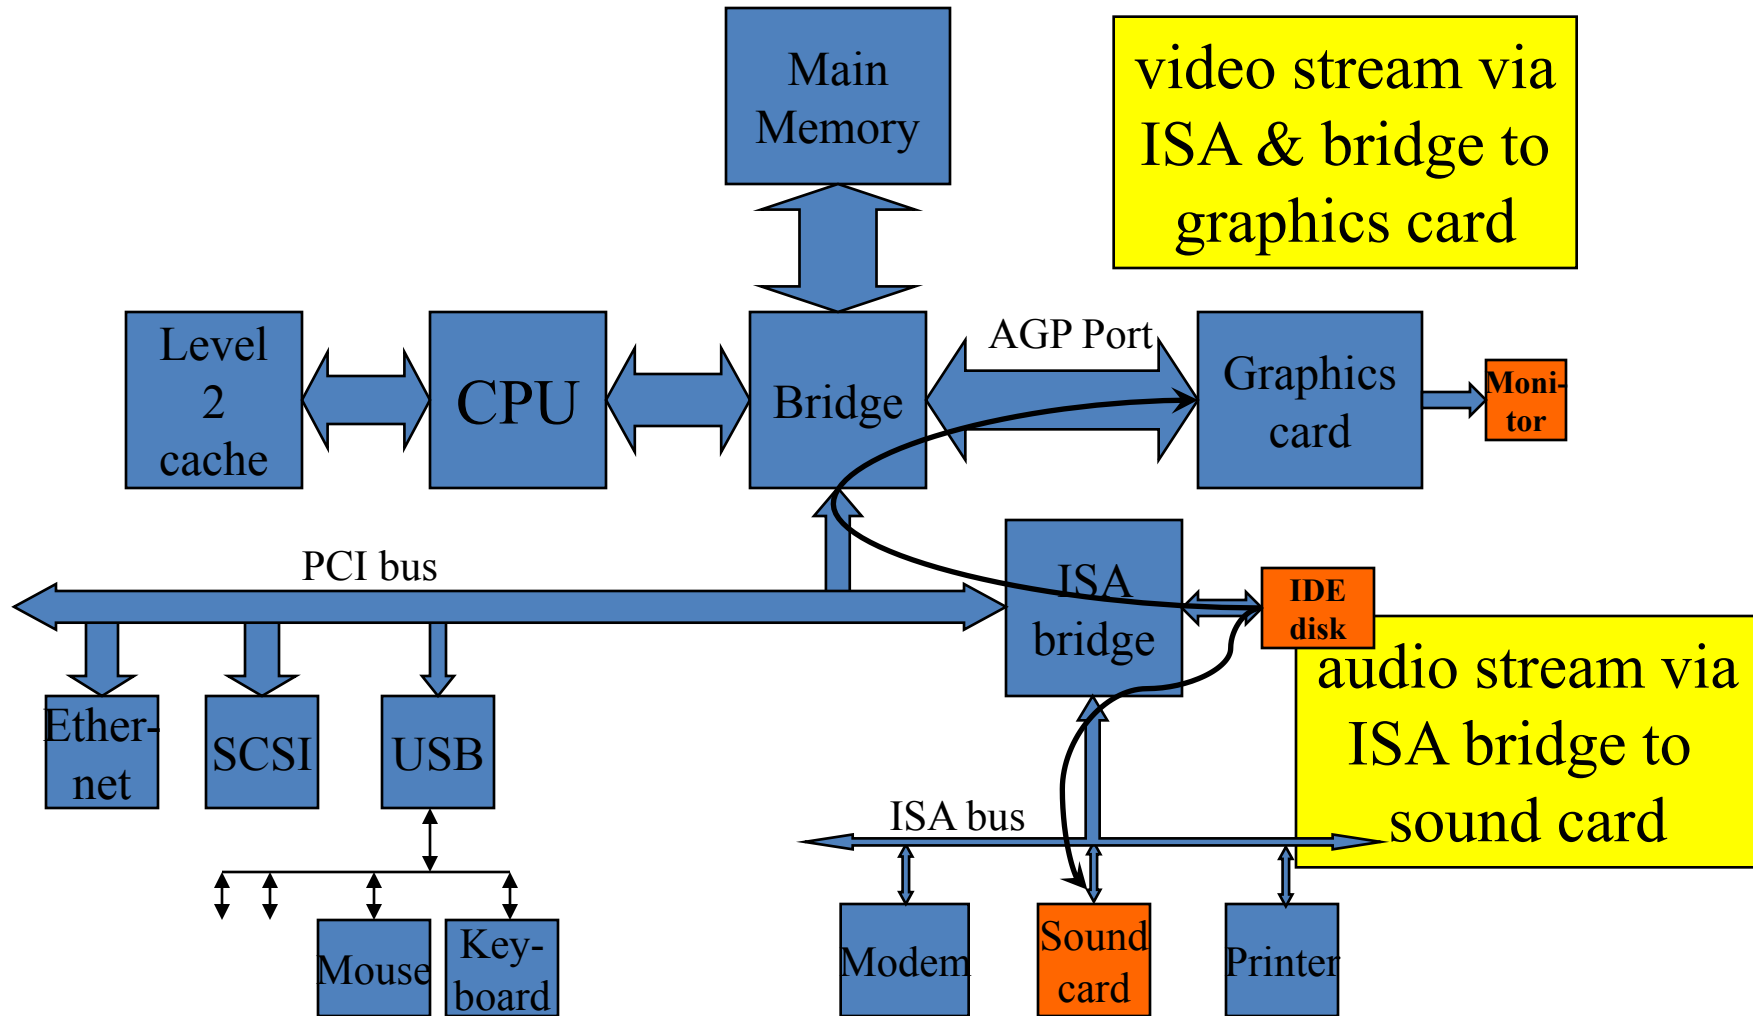


# System Organization (a decade ago)



- Separate data path for audio stream
  - Or headphone jack and volume control on CD drive itself
- Main system bus and CPU were too busy/slow to handle real-time audio

# System Organization (typical Pentium PC today)



# DVD devices

- Similar in concept, different in details
- More data, higher bandwidth
  - *Track spacing:* 0.834 $\mu$  vs. 1.6 $\mu$
  - *Bit spacing:* 0.4 $\mu$  vs. 0.74 $\mu$
  - *Capacity:* 4.7 Gbytes vs. 650 Mbytes

# Compression

An essential part of audio and video  
representation as files

# Why Compression? – CD-quality audio

- 22,050 Hz  $\Rightarrow$  44,100 samples/sec
  - 16 bits per sample
- Two channels  $\Rightarrow$  176,400 bytes/sec  
 $\cong$  1.4 mbits/sec
  - Okay for a modern PC
  - Not okay for 56 kb/sec modem (speed) or iPod (space)!
- MP-3  $\cong$  0.14 mbits/sec (10:1)
  - *Same* audio quality!
  - Compression ratio varies with type of music

# Why Compression? – Video

- “Standard” TV frame =  $640 \times 480$  pixels @ 25-30 frames/sec (fps)  $\Rightarrow$   
 $\Rightarrow 9,216,000$  pixels/sec =  $27,648,000$  bytes/sec
- Hollywood “standard” movie  $\leq 133$  minutes  
 $\Rightarrow$  approx. 210 gigabytes (standard resolution)!
- HDTV =  $1280 \times 720$  pixels @ 30 fps  
 $\Rightarrow 82,944,000$  bytes/sec (and rising!)
- DVD holds  $\sim 4.7$  gigabytes  
 $\Rightarrow$  average of  $\sim 620$  kilobytes/sec!
- “Standard” movie of 133 minutes requires serious compression just to *fit* onto DVD

# Video Compression Requirements

- Compression ratio  $> 50:1$ 
  - i.e., 210 gigabytes:4.7 gigabytes
- Visually indistinguishable from original
  - Even when paused
- Fast, cheap decoder
  - Slow encoder is okay
- VCR controls
  - Pause, fast forward, reverse

# Video Compression Standards

- MPEG (*Motion Picture Experts Group*)
  - Based on JPEG (*Joint Photographic Experts Group*)
  - Multi-layer
    - Layer 1 = system and timing information
    - Layer 2 = video stream
    - Layer 3 = audio and text streams
- Three standards
  - MPEG-1 – 352×240 frames; < 1.5 mb/sec ( < VHS quality)
    - Layer 3 = MP3 Audio standard
    - Typical uses:– video clips on internet; video for handhelds
  - MPEG-2 – standard TV & HDTV; 1.5-15 mb/sec
    - DVD encoding
  - MPEG-4 – combined audio, video, interactive graphics
    - 2D & 3D animations



# JPEG compression (single frame)

## 1. Convert RGB into YIQ

- $Y = \text{luminance}$  (i.e., brightness)  $\sim$  black-white TV
- $I, Q = \text{chrominance}$  (similar to *saturation* and *hue*)

Reason: Human eye is more sensitive to luminance than to color (rods vs. cones)


## 2. Down-sample $I, Q$ channels

- i.e., average over  $2 \times 2$  pixels to reduce resolution
- lossy compression, but barely noticeable to eye


## 3. Partition each channel into $8 \times 8$ blocks

- 4800  $Y$  blocks, 1200 each  $I$  &  $Q$  blocks (for  $640 \times 480$ )

# JPEG (continued)



# JPEG (continued)



# JPEG (continued)

4. Calculate *Discrete Cosine Transform* (DCT) of each  $8\times 8$  block
5. Divide  $8\times 8$  block of DCT values by  $8\times 8$  *quantization table*
  - Effectively throwing away higher frequencies
6. Linearize  $8\times 8$  block, run-length encode, and apply a Huffman code to reduce to a small fraction of original size (in bytes)


# JPEG (concluded)

7. Store or transmit  $8 \times 8$  quantization table followed by list of compressed blocks
- Achieves 20:1 compression with good visual characteristics
  - Higher compression ratios possible with visible degradation
- JPEG algorithm executed backwards to recover image
  - Visually indistinguishable from original @ 20:1
- JPEG algorithm is symmetric
  - Same speed forwards and backwards

# MPEG

- JPEG-like encoding of each frame
- Takes advantage of *temporal locality*
- I.e., each frame usually shares similarities with previous frame
  - ⇒ encode and transmit only differences
- Sometimes an object moves relative to background
  - ⇒ find object in previous frame, calculate difference, apply *motion vector*


# Temporal Locality (example)



Consecutive Video Frames

# MPEG organization

- Three types of frames
  - *I-frame: Intracoded or Independent.*
    - Full JPEG-encoded frame
    - Occurs at intervals of a second or so
    - Also at start of every *scene*
  - *P-frame: Predictive* frame
    - Difference from previous frame
  - *B-frame: Bidirectional* frame
    - Like *p-frame* but difference from both *previous* and *next* frame





# MPEG Characteristics

- Non-uniform data rate!
- Compression ratios of 50:1 – 80:1 are readily obtainable
- Asymmetric algorithm
  - Fast decode (like JPEG)
  - Encode requires image search and analysis to get high quality differences
- Cheap decoding chips available for
  - graphics cards
  - DVD players, etc.

# MPEG Problem – Fast Forward/Reverse

- Cannot simply skip frames
  - Next desired frame might be *B* or *P* derived from a skipped frame
- Options:
  - Separate fast forward and fast reverse files
    - MPEG encoding of every *n*th frame
    - Often used in video-on-demand server
  - Display only *I* and *P* frames
    - If *B* frame is needed, derive from nearest *I* or *P*

# “Movie” File Organization

- One MPEG-2 video stream
  - Possible fast forward, fast reverse sub-streams
- Multiple audio streams
  - Multiple languages
- Multiple text streams
  - Subtitles in multiple languages
- All interleaved
  - Possibly in multiple files!


# Frame

1 2 3 4 5 6 7 8

Video




English audio



French audio



German audio



English subtitles

Hello, Bob	Hello, Alice	Nice day	Sure is	How are you	Great	And you	Good
------------	--------------	----------	---------	-------------	-------	---------	------

Dutch subtitles

Dag, Bob	Dag, Alice	Mooie dag	Jazeker	Hoe gaat het	Prima	En jij	Goed
----------	------------	-----------	---------	--------------	-------	--------	------

Fast forward



Fast backward



# Operating System Challenge

- How to get the contents of a movie file from disk or DVD drive to video screen and speakers.
  - Fixed frame rate (25 or 30 fps)
  - Steady audio rate
  - Bounded *jitter*
- Such requirements are known as *Quality-of-Service* (QoS) guarantees.
- Classical problem in *real-time scheduling*
  - Obscure niche become mainstream!
  - See Silberschatz, §19.1–19.5

# QoS Guarantees

- Building a system to guarantee QoS effects the following:–
  - CPU processing
  - Scheduling and interrupt handling
  - File systems
  - Network protocols

# Requirement of Multimedia Operating Systems

- There are three levels of QoS
  - *Best-effort service*: the system makes a best effort with no attempt to guarantee
  - *Soft QoS*: allows different traffic streams to be prioritized, but no QoS guarantees are made
  - *Hard QoS*: system ensures QoS requirements are always met
    - Prioritization
    - Admission control
    - Bounded latency on interrupt handling, processing, etc.

# Parameters Defining QoS

- *Throughput*: the total amount of work completed during a specific time interval
- *Delay*: the elapsed time from when request is first submitted to desired result
- *Jitter*: delays that occur during playback of a stream.
- *Reliability*: how errors are handled during transmission and processing of continuous media



# Further QoS Issues

- QoS may be *negotiated* between the client and server.
- Operating systems may use an *admission control* algorithm
  - Admits a request for service only if the server has sufficient resources to satisfy the request

# Two common methods

- Rate Monotonic Scheduling
- Earliest Deadline First
- Many variations
- Many analytic methods for proving QoS (Quality of Service)


# Processor Scheduling for Real-Time

## *Rate Monotonic Scheduling (RMS)*

- Assume  $m$  periodic processes
  - Process  $i$  requires  $t_i$  msec of processing time every  $p_i$  msec.
  - Equal processing every interval — like clockwork!

# Example

- Periodic process  $i$  requires the CPU at specified intervals (periods)
- $p_i$  is the duration of the period
- $t_i$  is the processing time
- $d_i$  is the deadline by when the process must be serviced
  - Often same as end of period



# Processor Scheduling for Real-Time

## *Rate Monotonic Scheduling (RMS)*

- Assume  $m$  periodic processes
  - Process  $i$  requires  $t_i$  msec of processing time every  $p_i$  msec.
  - Equal processing every interval — like clockwork!
- Assume
$$\sum_{i=1}^m \frac{t_i}{p_i} \leq 1$$
- Assign priority of process  $i$  to be
  - Statically assigned
- Let priority of non-real-time processes be 0


$$\frac{1}{p_i}$$

## Rate Monotonic Scheduling (continued)

- Scheduler simply runs highest priority process that is ready
  - May pre-empt other real-time processes
  - Real-time processes become ready in time for each frame or sound interval
  - Non-real-time processes run only when no real-time process needs CPU


# Example

- $p_1 = 50$  msec;  $t_1 = 20$  msec
- $p_2 = 100$  msec;  $t_2 = 35$  msec
- $\text{Priority}(p_1) > \text{Priority}(p_2)$
- Total compute load is 75 msec per every 100 msec.
- Both tasks complete within every period
  - 25 msec per 100 msec to spare



# Example 2

- $p_1 = 50$  msec;  $t_1 = 25$  msec
- $p_2 = 80$  msec;  $t_2 = 35$  msec
- $\text{Priority}(p_1) > \text{Priority}(p_2)$
- Total compute load is  $\sim 94\%$  of CPU.
- Cannot complete both tasks within some periods
  - Even though there is still CPU capacity to spare!





# Processor Scheduling for Real-Time

## *Earliest Deadline First (EDF)*

- When each process  $i$  become ready, it announces deadline  $D_i$  for its next *task*.
- Scheduler always assigns processor to process with earliest deadline.
  - May pre-empt other real-time processes

# Earliest Deadline First Scheduling (continued)


- No assumption of periodicity
- No assumption of uniform processing times
- **Theorem:** If *any* scheduling policy can satisfy QoS requirement for a sequence of real time tasks, then EDF can also satisfy it.
  - *Proof:* If  $i$  scheduled before  $i+1$ , but  $D_{i+1} < D_i$ , then  $i$  and  $i+1$  can be interchanged without affecting QoS guarantee to either one.

## Earliest Deadline First Scheduling (continued)

- EDF is more complex scheduling algorithm
  - Priorities are dynamically calculated
  - Processes must know deadlines for tasks
- EDF can make higher use of processor than RMS
  - Up to 100%
- There is a large body of knowledge and theorems about EDF analysis

# Example 2 (again)

- Priorities are assigned according to deadlines:
  - the earlier the deadline, the higher the priority;
  - the later the deadline, the lower the priority.



# Network Management (continued)

- Broadcasting
  - Like cable TV
  - Fixed portion of network bandwidth dedicated to set of broadcast streams; inflexible
- Multicasting
  - Typical webcast
  - Multicast tree set up through internet to reach customers
  - Customers can come and go
- Unicast
  - Dedicated connection between server and client
  - Streaming version of familiar transport protocols

# Streaming – Delivery of Multimedia Data over Network

- Two types of streaming:–
  - *Progressive download*: client begins playback of multimedia file during delivery.
    - File is ultimately stored on client computer
    - (Hopefully) download speed > playback speed
  - *Real-time streaming*: multimedia file is delivered to, but *not* stored on, client computer
    - Played back at same speed as delivery
    - Limited amount of buffering to remove jitter

# Two Types of Real-time Streaming


- *Live streaming*: to deliver a live event while it is occurring.
  - Broadcast or multicast
- *On-demand streaming*: to deliver archived media streams
  - Movies, lectures, old TV shows, etc.
  - Events *not* delivered when they occur.
  - Playback with *pause, fast forward, reverse, etc*
  - Unicast (usually)

# Network Streaming

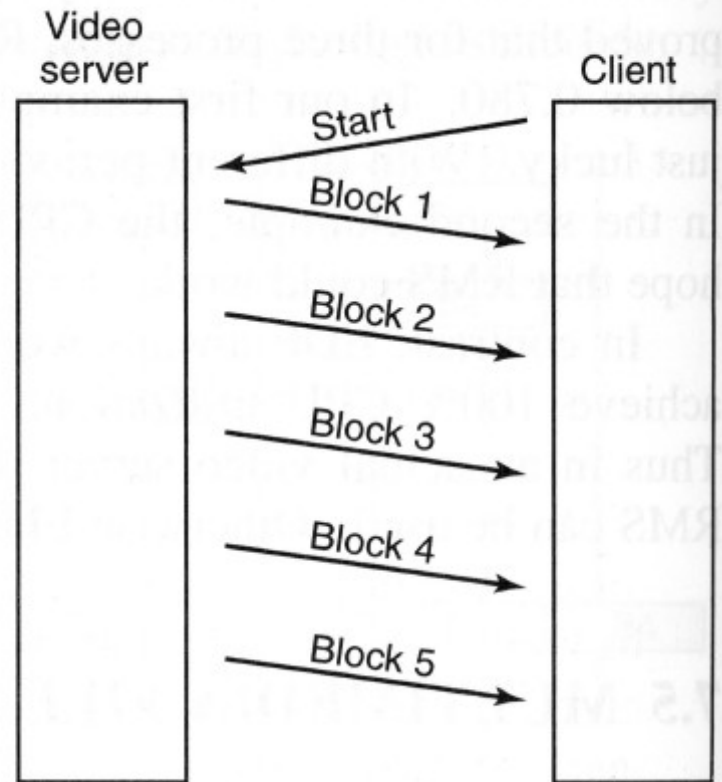
- Traditional HTTP
  - Stateless
  - Server responds to each request independently
- Real-Time Streaming Protocol (RTSP)
  - Client initiates a “push” request for stream
  - Server provides media stream at frame rate



# Pull (HTTP) vs. Push (RTSP) server



(a)

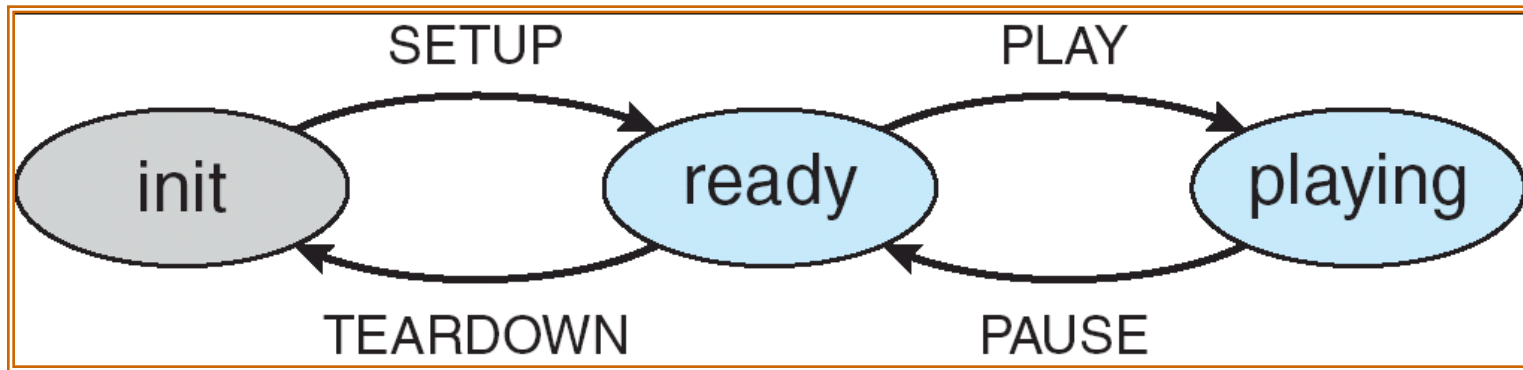


(b)

# RTSP States

- *SETUP*: server allocates resources for client session.
- *PLAY*: server delivers stream to a client session.
- *PAUSE*: server suspends delivery of a stream.
- *TEARDOWN*: server releases resources and breaks down connection.

# RTSP state machine



# Bandwidth Negotiation


- Client application provides feedback to server to adjust bandwidth
- E.g.,
  - Windows Media Player
  - RealPlayer
  - Quicktime

# Video Server

- Multiple CPUs
- Disk farm
  - 1000s of disks
- Multiple high-bandwidth network links
  - Cable TV
  - Video on demand
  - Internet

# Multimedia File & Disk Management



- Single movie or multimedia file on disk
  - Interleave audio, video, etc.
    - So temporally equivalent blocks are near each other
  - Attempt contiguous allocation
    - Avoid seeks within a frame



# File organization – Frame vs. Block

- Frame organization
  - Small disk blocks (4-16 Kbytes)
  - Frame index entries point to starting block for each frame
  - Frames vary in size (MPEG)
  - Advantage: very little storage fragmentation
  - Disadvantage: large frame table in RAM
- Block organization
  - Large disk block (256 Kbytes or more)
  - Block index entries point to first *I-frame* of a sequence
  - Multiple frames per block
  - Advantage: much smaller block table in RAM
  - Disadvantage: large storage fragmentation on disk

# Frame vs. Block organization






# File Placement on Server

- Random
- Striped
- “Organ pipe” allocation
  - Most popular video in center of disk
  - Next most popular on either side of it, etc.
  - Least popular at edges of disk
  - Minimizes seek distance

# Resources on a file server



# Conclusion

- Multimedia computing is challenging
- Possible with modern computers
  - Compression is essential, especially for video
- Real-time computing techniques move into mainstream
  - Processor and disk scheduling
- There is much more to this subject than fits into one class

Thank You.....  
For your attentions